

Inference latency: QUEST-KG is 19× to 500× faster than LLM baselines

